

# ***COMP 7990***

## ***Principles and Practices of Data Analytics***

*Lecture 5: Artificial Neural Network,  
Support Vector Machine and  $k$  Nearest  
Neighbors ( $k$ -NN) Algorithm*

*Dr. Eric Lu Zhang*

# Outline for Data Preprocessing and Data Mining

- Data Preprocessing
- **Supervised learning**

- ❖ Regression

1. Linear regression with one variable
2. Linear Regression with multiple variables
3. The relationship between Correlation and Regression

- ❖ **Classification**

1. Perceptron
2. Artificial Neural Network
3. Support Vector Machine
4. K Nearest Neighbor

- **Unsupervised learning**

1. K-means Clustering
2. Hierarchical Clustering

# Outline for Data Preprocessing and Data Mining

- Data Preprocessing
- **Supervised learning**

- ❖ Regression

1. Linear regression with one variable
2. Linear Regression with multiple variables
3. The relationship between Correlation and Regression


- ❖ **Classification**

1. Perceptron
2. Artificial Neural Network
3. Support Vector Machine
4. K Nearest Neighbor

- **Unsupervised learning**


1. K-means Clustering
2. Hierarchical Clustering

# Artificial Neural Networks




# Recall on Perceptron

- A linear classification function  $f(\mathbf{x}) = \sum_{j=1}^n w_j x_j + b$ 
  - $y = 1$  if  $f(\mathbf{x}) = \sum_{j=1}^n w_j x_j + b > 0$
  - $y = -1$  if  $f(\mathbf{x}) = \sum_{j=1}^n w_j x_j + b < 0$






- $[x_1, \dots, x_n]$ : inputs
- $[w_1, \dots, w_n]$ : weights
- $b$ : bias term
- $\Sigma$ : summation
- $f$ : activation function (sign function used)

# Simplified Illustration of a Neuron





# Limitation of Perceptron



- Many classification problems are NOT linearly separable.



- Possible solution for nonlinear classification problem: Composition
  - Multiple Layer Perceptron (also called Feedforward Neural Network)




Compose  
them together






# Multi-Layer Perceptron for Nonlinear Classification



$x_1$	$x_2$	$y$
1	-1	1
1	1	-1
-1	1	1
-1	-1	-1




This multi-layer perceptron can solve this nonlinear classification problem.

How to learn the model parameter (i.e., weights on the edge) from data?

# Multi-Layer Perceptron

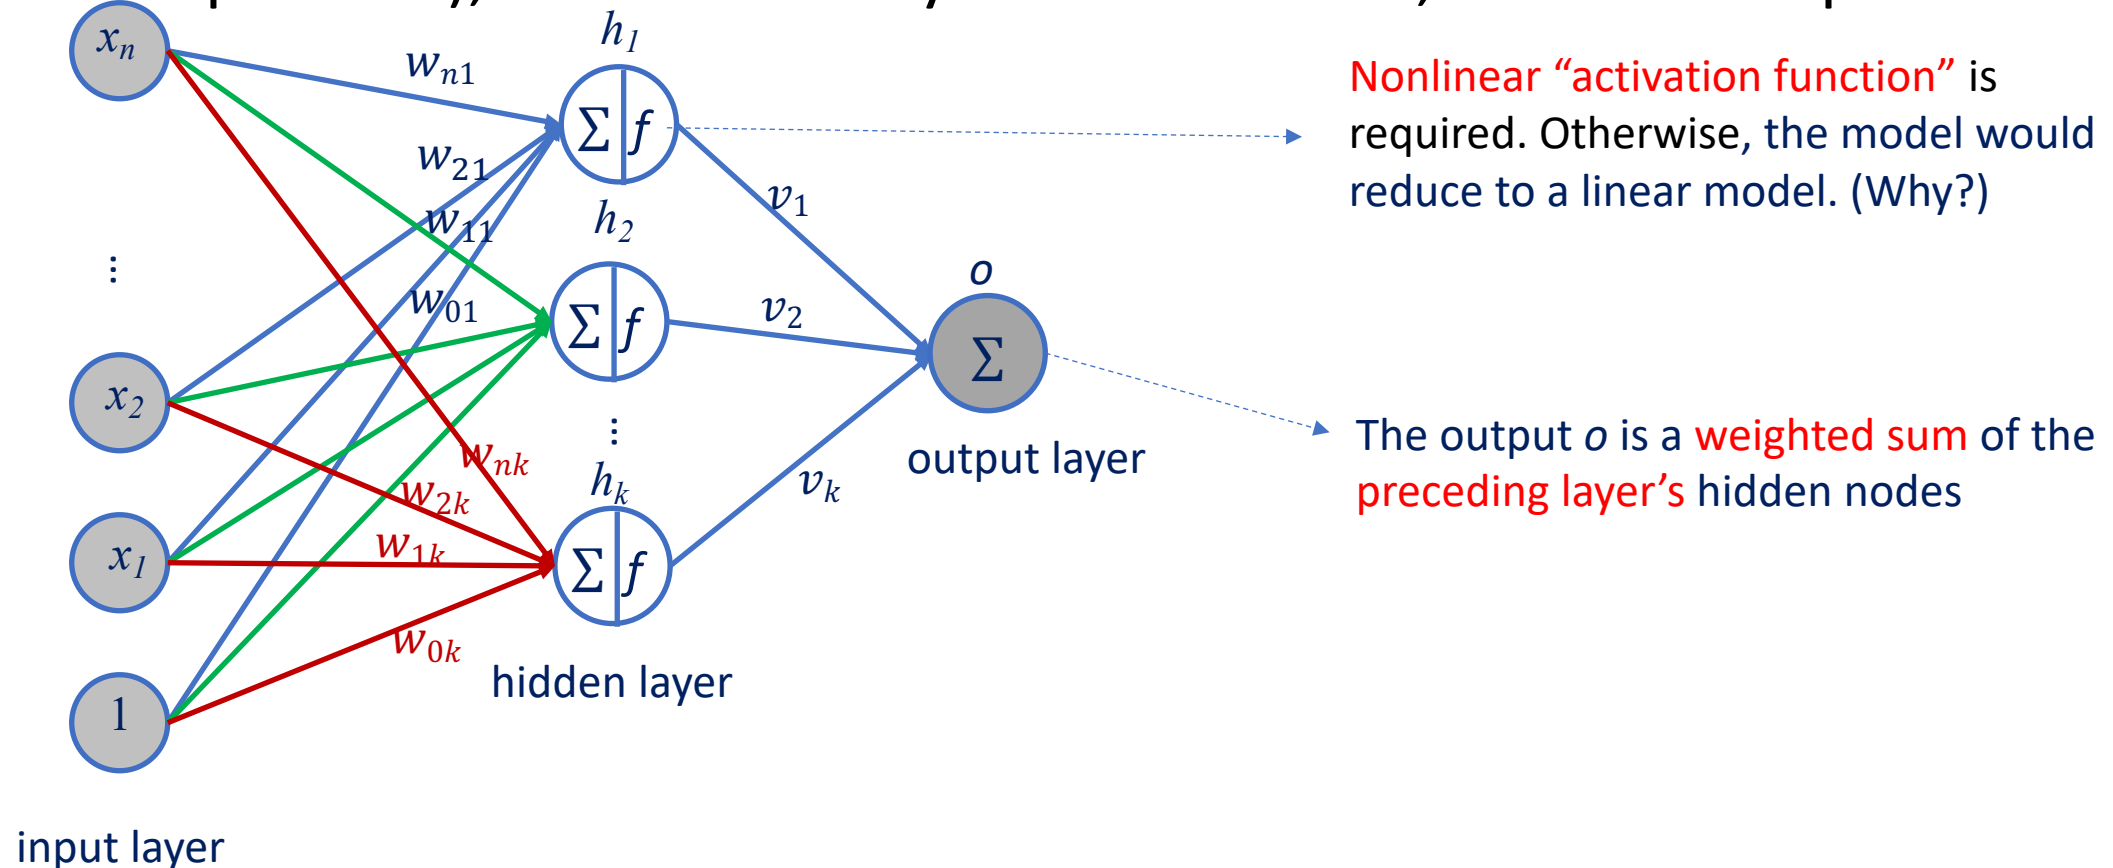
- Composed of several Perceptron-like units arranged in multiple layers




- Consists of an **input layer**, one or more **hidden layers**, and an output layer
- Nodes in the hidden layers carry out nonlinear transformation of the inputs
- Universal Approximator (Hornik, 1991)

# Multi-Layer Perceptron (MLP) with One Hidden Layer

- Multi-Layer Perceptron with  $n$  inputs (i.e. the dimension of input samples is  $n$ ), one hidden layer with  $k$  nodes, and one output.



# Output of MLP with 2 Hidden Nodes and 2-dimensional Input



$$o = v_1 h_1 + v_2 h_2 = \sum_{i=1}^2 v_i h_i$$

$$h_1 = f\left(\sum_{i=0}^2 w_{i1} x_i\right) = f(\mathbf{w}_1^T \mathbf{x})$$


$$h_2 = f\left(\sum_{i=0}^2 \mathbf{w}_{i2} x_i\right) = f(\mathbf{w}_2^T \mathbf{x})$$

$$W = \begin{bmatrix} w_{0,1} & w_{0,2} \\ w_{1,1} & w_{1,2} \\ w_{2,1} & w_{2,2} \end{bmatrix} = [\mathbf{w}_1 \quad \mathbf{w}_2]$$

$$W^T = \begin{bmatrix} w_{0,1} & w_{1,1} & w_{2,1} \\ w_{0,2} & w_{1,2} & w_{2,2} \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^T \\ \mathbf{w}_2^T \end{bmatrix}$$

$W$  is a 3-by-2 matrix with the  $j$ -th column in  $W$  denoting the weights of the  $j$ -th node in the hidden layer.

# Output of MLP with $k$ Hidden Nodes and $d$ -dimensional Input



$W$  is a  $(n+1)*k$  matrix, the  $j$ -th column in  $W$  denotes the weights of the  $j$ -th node in the hidden layer.

$$o = \sum_{i=1}^k v_i h_i = \mathbf{v}^T \mathbf{h}$$


$$\mathbf{h} = f(W^T \mathbf{x})$$

$$\begin{aligned} h_1 &= f\left(\sum_{i=0}^n w_{i,1} x_i\right) = f(\mathbf{w}_1^T \mathbf{x}) \\ h_2 &= f\left(\sum_{i=0}^n w_{i,2} x_i\right) = f(\mathbf{w}_2^T \mathbf{x}) \\ &\vdots \\ h_k &= f\left(\sum_{i=0}^n w_{i,k} x_i\right) = f(\mathbf{w}_k^T \mathbf{x}) \end{aligned}$$


$$o = \mathbf{v}^T f(W^T \mathbf{x})$$

# Commonly Used Nonlinear Activation Functions

- Some common choices for nonlinear activation function  $f$ 
  - Sigmoid:  $f(x) = \sigma(x) = \frac{1}{1+\exp(-x)}$  (range between 0-1)
  - Tanh:  $f(x) = 2\sigma(2x) - 1 = \frac{2}{1+\exp(-2x)} - 1$  (range between -1 and +1)
  - Rectified Linear Unit (ReLU):  $f(x) = \max(0, x)$



# Multi-Layer Perceptron



- For an MLP with  $L$  hidden layers,  $\mathbf{h}^{(1)}, \dots, \mathbf{h}^{(L)}$  and the scalar-valued output is computed as

$$o = \sum_{i=1}^k h_i v_i = \mathbf{v}^T \mathbf{h}^{(L)}$$

$$\mathbf{h}^{(L)} = f(W^{(L)T} \mathbf{h}^{(L-1)})$$

$$\mathbf{h}^{(L-1)} = f(W^{(L-1)T} \mathbf{h}^{(L-2)})$$

$$\vdots$$


$$\mathbf{h}^{(2)} = f(W^{(2)T} \mathbf{h}^{(1)})$$

$$\mathbf{h}^{(1)} = f(W^{(1)T} \mathbf{x})$$

Why nonlinear “activation function” is required?

# Why nonlinear “activation function” is required


- If we remove the nonlinear “activation function”, the output of the neural network will be reduced to

$$\begin{aligned} o &= \sum_{i=1}^k h_i v_i = \mathbf{v}^T \mathbf{h}^{(L)} \\ &= \mathbf{v}^T W^{(L)T} \mathbf{h}^{(L-1)} \\ &= \mathbf{v}^T W^{(L)T} W^{(L-1)T} \mathbf{h}^{(L-2)} \\ &\dots \\ &= \mathbf{v}^T W^{(L)T} W^{(L-1)T} \dots W^{(2)T} W^{(1)T} \mathbf{x} \end{aligned}$$


All these operations are linear transformation.  
Therefore, without nonlinear activation  
function, it will reduce to a simple linear  
model.




# Learning MLP Via Backpropagation



- For a given training dataset, we want to learn the parameter  $(\mathbf{W}^{(1)}, \dots, \mathbf{W}^{(L)}, \mathbf{v})$  by minimizing some loss function.
- **Backpropagation** (gradient descent + chain rule for derivatives) is commonly used to do this efficiently.

# Learning MLP (one hidden layer)



- Given one data sample of input and true output  $\{\mathbf{x}, y\}$ , our goal is to minimize  $(y - o)^2 = (y - \mathbf{v}^T f(W^T \mathbf{x}))^2$
- Given  $m$  data sample of input and true output  $\{\{\mathbf{x}_1, y_1\}, \{\mathbf{x}_2, y_2\}, \dots, \{\mathbf{x}_m, y_m\}\}$ , the goal becomes:

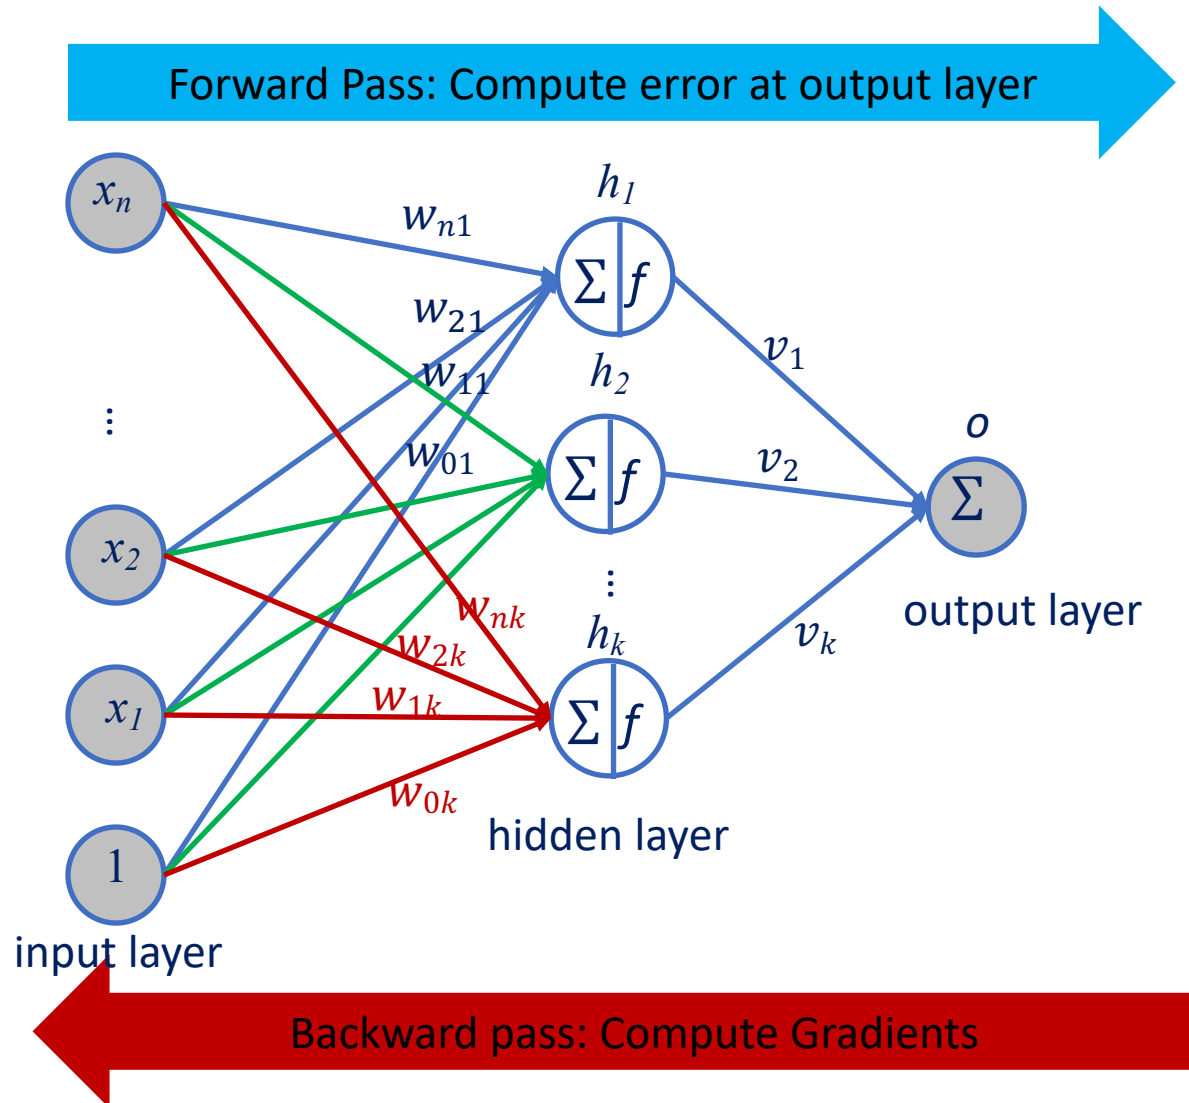
$$\min_{W, \mathbf{v}} \frac{1}{2m} \sum_{i=1}^m (y_i - o_i)^2$$

$$\min_{W, \mathbf{v}} \frac{1}{2m} \sum_{i=1}^m (y_i - \mathbf{v}^T f(W^T \mathbf{x}_i))^2$$

$$= \min_{W, \mathbf{v}} \frac{1}{2m} \sum_{i=1}^m \left( y_i - \sum_{j=1}^k v_j f(\mathbf{w}_j^T \mathbf{x}_i) \right)^2$$

*Note: We use square loss here for simply illustrating the key idea of learning MLP. In practice, we usually use hinge loss or cross entropy loss (log loss) for classification problem.*


# Backpropagation



- Basically consists of a forward pass and a backward pass
- Forward pass computes the error  $e_i$  using the current parameters
- Backwards pass computes the gradient and updates the parameters, starting from the parameter at the rightmost layer and the moving backwards.
- Also good at reusing previous computations (updates of parameters at any layer depend on parameters of the upper layer)


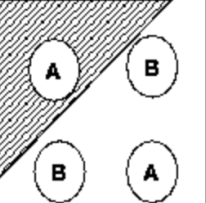
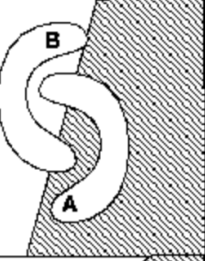

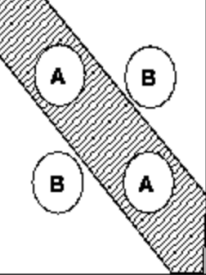
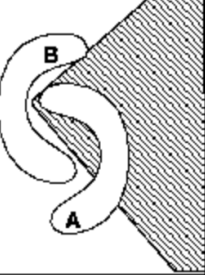

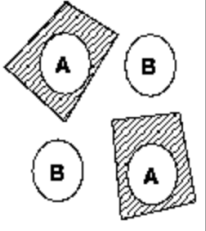
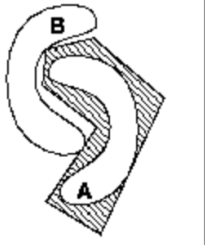
# Solutions are Local Minima

*Parameter Initialization and Learning Rate Do Matter*






# Pros and Cons of MLP

- Pros
  - Versatile: Adaptive to many datasets
  - Can capture nonlinear dependence of input and output
- Cons
  - Does not work for small data sets
  - Blackbox; Hard to interpret
  - Speed of convergence
  - Local minimum
  - Overfitting issue (how to select the structure; how to achieve good generalization)

Structure	Regions	XOR	Meshed regions
single layer 	Half plane bounded by hyper-plane		
two layer 	Convex open or closed regions		
three layer 	Arbitrary (limited by # of nodes)		

# Underfitting/Overfitting and Model Complexity



# Outline for Data Preprocessing and Data Mining

- Data Preprocessing
- **Supervised learning**

- ❖ Regression

1. Linear regression with one variable
2. Linear Regression with multiple variables
3. The relationship between Correlation and Regression

- ❖ **Classification**


1. Perceptron
2. Artificial Neural Network
3. Support Vector Machine
4. K Nearest Neighbor

- **Unsupervised learning**

1. K-means Clustering
2. Hierarchical Clustering

# What Is the Best Hyperplane Separator?

- Perceptron finds one of many possible hyperplanes separating the data
  - If the hyperplane exist
- Among the many possible hyperplanes, which one is the best?




- Intuitively, we want the hyperplane not **too** close to each of the classes. In other words, the one **with** the maximum margin is preferred.
- A large margin can lead to good generalization on the test (unseen) data.




# Support Vector Machine (SVM)


- Probably the most popular/influential classification algorithm
- Backed by solid theoretical groundings
- A hyperplane based classifier (like the Perceptron)
- Additionally uses the maximum margin Principle
  - Finds the hyperplane with maximum separation margin on the training data



# Calculate distance from a point to a line



# The Concept of Margins



**Margin of a sample :**  $\gamma_i$  of a sample  $\mathbf{x}_i$  is its distance from the hyperplane

$$\|\mathbf{w}\| = \sqrt{\sum_{i=1}^n w_i^2} \quad \|\mathbf{w}\|^2 = \sum_{i=1}^n w_i^2$$

$$\gamma_i = \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|}$$


**Margin of a set:** is the minimum margin of all samples

$$\gamma = \min_{1 \leq i \leq m} \gamma_i = \min_{1 \leq i \leq m} \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|}$$

If we are asked to move the hyperplane to achieve better generalization, what should we do?

# Support Vector Machine

- A hyperplane based linear classifier defined by  $\mathbf{w}$  and  $b$
- Prediction rule:  $y = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$
- **Given:** Training data  $\{\mathbf{x}_i, y_i\}_{i=1}^m$
- **Goal:** Learn  $\mathbf{w}$  and  $b$  that achieve the maximum margin
- For now, assume the entire training data is linearly separable. We will handle linearly unseparable cases later




- Assume the hyperplane is such that

- $\mathbf{w}^T \mathbf{x}_i + b \geq 1$  for  $y_i = +1$
- $\mathbf{w}^T \mathbf{x}_i + b \leq -1$  for  $y_i = -1$
- Equivalently,  $y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1$
- The hyperplane's margin:

$$\gamma = \min_{1 \leq i \leq m} \gamma_i = \min_{1 \leq i \leq m} \frac{|\mathbf{w}^T \mathbf{x}_i + b|}{\|\mathbf{w}\|} = \frac{1}{\|\mathbf{w}\|}$$

An orange curved arrow points from the expression  $\min_{1 \leq i \leq m} |\mathbf{w}^T \mathbf{x}_i + b| = 1$  to the numerator of the margin formula.




# Support Vector Machine: the optimization problem

- We want to maximize the margin  $\gamma = \frac{1}{\|\mathbf{w}\|}$ , which is equivalent to minimize  $\|\mathbf{w}\|$  or  $\frac{\|\mathbf{w}\|^2}{2}$
- Therefore, the optimization problem of SVM for the *separable case* would be

$$\begin{aligned} \min \quad & \frac{\|\mathbf{w}\|^2}{2} \\ \text{subject to} \quad & y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, i = 1, \dots, m \end{aligned}$$

This is a Quadratic Program (QP) with  $n$  linear inequality constraints.



# SVM: Solving the Optimization Problem (optional)

- The optimization problem is

$$\begin{aligned} \min \quad & \frac{\|\mathbf{w}\|^2}{2} \\ \text{subject to} \quad & y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, i = 1, \dots, m \end{aligned}$$

- Introducing Lagrange Multipliers  $\alpha_i$ , one for each constraint, leads to the primal Lagrangian:

$$\begin{aligned} \min L_p = \quad & \frac{\|\mathbf{w}\|^2}{2} + \sum_{i=1}^m \alpha_i (1 - y_i(\mathbf{w}^T \mathbf{x}_i + b)) \\ \text{subject to} \quad & \alpha_i \geq 0, i = 1, \dots, m \end{aligned}$$

# SVM: Solution! (optional)


- Once we have the  $\alpha_i$ , we can compute  $\mathbf{w}$  and  $b$  as:

(duality optimization, KKT, Sequential Minimal Optimization)

$$\mathbf{w} = \sum_{i=1}^m \alpha_i y_i \mathbf{x}_i \quad b = y_i - \sum_{j=1}^m \alpha_j y_j \mathbf{x}_j^T \mathbf{x}_i$$

$$\mathbf{w}^T \mathbf{x}_i + b = y_i$$

*any support vector  
will work*




- An important consequence:
  - $\alpha_i$  is non-zero only if  $\mathbf{x}_i$  lies on one of the two margin boundaries, i.e., for which
$$y_i(\mathbf{w}^T \mathbf{x}_i + b) = 1$$
  - The samples are called **support vectors**.



# SVM: Non-separable case

- Non-separable case: No hyperplane can separate the classes perfectly (common in practice).
- Still want to find the maximum margin hyperplane, but...
  - Allow some training samples to be misclassified (can you identify those points on the right?)
  - Allow some training samples to fall within the margin region (can you identify those points on the right?)



# SVM: Use Slack Variables

- Solution: introduce slack variables
- Recall: for the separable case, the constraints are:

$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1, i = 1, \dots, m$$

- For the non-separable case, we relax the constraints by adding slack variables


$$y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, i = 1, \dots, m$$

- $\xi_i \geq 0$  is required
- $\xi_i > 1$  for misclassified samples (for outside the margin)

# Support Vectors for non-separable cases

- Recall: the separable case has only one type of support vectors
  - Ones that lies on the margin boundaries  $\mathbf{w}^T \mathbf{x}_i + b = 1$  or  $\mathbf{w}^T \mathbf{x}_i + b = -1$
- The non-separable case has three types of support vectors
  1. Lying on the margin boundaries  $\xi_i = 0$
  2. Lying with the margin region  $0 < \xi_i < 1$  but still on the correct side
  3. Lying on the wrong side the hyperplane  $\xi_i \geq 1$

*Can you identify them on the right?*



# The Optimization Problem (optional)

- While we “allow” misclassified training samples
  - We want the number of misclassified training samples to be minimized
  - By minimizing the sum of slack variables  $\sum_{i=1}^m \xi_i$
- Therefore, the new optimization problem for non-separable case will be

$$\begin{aligned} & \min \frac{\|\mathbf{w}\|^2}{2} + C \sum_{i=1}^m \xi_i \\ & \text{subject to } y_i(\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i, i = 1, \dots, m \\ & \quad \xi_i \geq 0 \end{aligned}$$

$C$  is a hyper-parameter to control the tradeoff between **training errors** and **margins**

- Large  $C$ , prefer low **training errors**
- Small  $C$ , prefer large **margins**

# The Optimization Problem (optional)


- As in the linearly separable problem, by following the derivations, we will obtain the following dual problem

$$\begin{aligned} \max L_d &= \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j (\mathbf{x}_i^T \mathbf{x}_j) \\ \text{subject to } &\sum_{i=1}^m \alpha_i y_i = 0, 0 \leq \alpha_i \leq C, i = 1, \dots, m \end{aligned}$$

- Again a Quadratic Programming problem for  $\alpha$
- Given  $\alpha$ , the solution  $\mathbf{w}$ ,  $b$  has the same form as the separable case
- Note:  $\alpha$  is again sparse. Non-zero  $\alpha_i$  corresponds to the **support vectors**

# SVM for Nonlinear Classification



- Problem: SVM with linear function  $\mathbf{w}^T \mathbf{x} + b$  have very limited representation power. Therefore, it can not solve nonlinear classification problem.



- Good news: With a slight modification using **kernel trick**, SVM can solve highly nonlinear classification problems.

# Kernel SVM for Nonlinear Classification


- Key idea: Projecting the input to a high dimensional feature space so that non-linear classification problem becomes linearly separable again!



# Kernels


- Kernels: Make linear models work in nonlinear settings
  - By mapping data to higher dimensions where it exhibits linear patterns.
  - Apply the linear model in the new input space
  - Mapping means changing the feature representation

- Examples:







# Kernels



- Kernels: Make linear models work in nonlinear settings
  - By mapping data to higher dimensions where it exhibits linear patterns.
  - Apply the linear model in the new input space
  - Mapping means changing the feature representation
  - <https://www.youtube.com/watch?v=3liCbRZPrZA>
- Examples:



$$\mathbf{x}: [x_1, x_2] \rightarrow \mathbf{z}: [x_1^2, \sqrt{2}x_1x_2, x_2^2]$$



# Feature Mapping (optional)

- Consider the following mapping  $\phi$  for a sample  $\mathbf{x} = [x_1, x_2, \dots, x_n]$

$$\phi: \mathbf{x} \rightarrow [x_1^2, x_2^2, \dots, x_n^2, x_1x_2, x_1x_3, \dots, x_1x_n, \dots, x_{n-1}x_n]$$

- It is an example of quadratic mapping
  - Each new feature uses a pair of the original features
- Problem: Explicit mapping leads to the number of features blow up!
- Fortunately, kernel trick help us to avoid the problem.
  - The mapping does not have to be explicitly computed

# Kernel as high dimensional feature mapping (optional)

- Consider two samples  $\mathbf{x}$ :  $[x_1, x_2]$  and  $\mathbf{z}$ :  $[z_1, z_2]$
- Let us assume there is a kernel function  $k$  that takes inputs  $\mathbf{x}$  and  $\mathbf{z}$

$$\begin{aligned}k(\mathbf{x}, \mathbf{z}) &= (\mathbf{xz})^2 \\&= (x_1z_1 + x_2z_2)^2 \\&= x_1^2z_1^2 + x_2^2z_2^2 + 2x_1x_2z_1z_2 \\&= (x_1^2, \sqrt{2}x_1x_2, x_2^2)(z_1^2, \sqrt{2}z_1z_2, z_2^2) \\&= \phi(\mathbf{x})\phi(\mathbf{z})\end{aligned}$$

$\phi(\mathbf{x})\phi(\mathbf{z})$  is computed efficiently in original input space.

- This kernel function  $k$  implicitly defines a mapping  $\phi$  to a higher dimensional space

$$\phi(\mathbf{x}) = [x_1^2, \sqrt{2}x_1x_2, x_2^2]$$

- Note that we do not have to define/compute this mapping. Simply defining the kernel is a certain way to give a higher dimensional mapping  $\phi$

# Kernel: Formal definition (optional)

- $\phi$  takes input  $\mathbf{x}$  in input space and maps to feature space
- Kernel  $k(\mathbf{x}, \mathbf{z})$  takes two inputs and gives their similarity in feature space

$$k(\mathbf{x}, \mathbf{z}) = \phi(\mathbf{x})\phi(\mathbf{z})$$

- Some Examples of Kernels
  - Linear kernel  $k(\mathbf{x}, \mathbf{z}) = \mathbf{xz}$
  - Quadratic Kernel  $k(\mathbf{x}, \mathbf{z}) = (\mathbf{xz})^2$  or  $k(\mathbf{x}, \mathbf{z}) = (1 + \mathbf{xz})^2$
  - Polynomial Kernel  $k(\mathbf{x}, \mathbf{z}) = (\mathbf{xz})^q$  or  $k(\mathbf{x}, \mathbf{z}) = (1 + \mathbf{xz})^q$
  - Radial Basis Function (RBF) kernel  $k(\mathbf{x}, \mathbf{z}) = \exp(-\gamma \|\mathbf{x} - \mathbf{z}\|^2)$ 
    - The RBF kernel corresponds to an infinite dimensional feature space. We can not actually write down  $\phi(\mathbf{x})$  for RBF kernel.

# Using Kernel

- Kernel can turn a linear model into a nonlinear one
- Recall: Kernel  $k(\mathbf{x}, \mathbf{z})$  represents a dot product in some high dimensional feature space.
- Any learning algorithm in which examples only appear as dot products  $(\mathbf{x}_i \mathbf{x}_j)$  can be kernelized (i.e., non-linearized)
  - by replacing the  $(\mathbf{x}_i \mathbf{x}_j)$  by  $\phi(\mathbf{x}_i) \phi(\mathbf{x}_j) = k(\mathbf{x}_i, \mathbf{x}_j)$
- Most learning algorithms can be kernelized:
  - Perceptron, SVM, linear regression, logistic regression, etc

# Kernelize SVM Training (optional)

- Recall the dual Lagrangian for linear SVM

$$\begin{aligned}\max L_d &= \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j (\mathbf{x}_i \mathbf{x}_j) \\ \text{subject to } &\sum_{i=1}^m \alpha_i y_i = 0, 0 \leq \alpha_i \leq C, i = 1, \dots, m\end{aligned}$$

- Replace  $(\mathbf{x}_i \mathbf{x}_j)$  by  $\phi(\mathbf{x}_i) \phi(\mathbf{x}_j) = k(\mathbf{x}_i, \mathbf{x}_j)$ , where  $k(,)$  is some kernel function

$$\begin{aligned}\max L_d &= \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j k(\mathbf{x}_i, \mathbf{x}_j) \\ \text{subject to } &\sum_{i=1}^m \alpha_i y_i = 0, 0 \leq \alpha_i \leq C, i = 1, \dots, m\end{aligned}$$

- Now, SVM learns a linear separator in kernel defined feature space
  - This corresponds to non-linear separator in the original input space

# Kernel SVM for Nonlinear Classification (optional)

- For a new input  $\mathbf{x}$ , the output will be:

$$y = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^m \alpha_i y_i \mathbf{x}_i^T \mathbf{x} \quad \text{Since } \mathbf{w} = \sum_{i=1}^m \alpha_i y_i \mathbf{x}_i$$

- For the non-linear classification, apply non-linear transformation  $\Phi(\mathbf{x})$  to project  $\mathbf{x}$  to a higher dimensional space and the inner product term becomes

$$y = \sum_{i=1}^m \alpha_i y_i \Phi(\mathbf{x}_i)^T \Phi(\mathbf{x})$$

# Kernel SVM for Nonlinear Classification (optional)

- As only the inner product is needed, we can apply the kernel trick. That is we care only the the way to measure distance between two points.

$$y = \sum_{i=1}^m \alpha_i y_i \Phi(\mathbf{x}_i)^T \Phi(\mathbf{x})$$

$$y = \mathbf{w}^T \mathbf{x} = \sum_{i=1}^m \alpha_i y_i k(\mathbf{x}_i, \mathbf{x})$$

Note: We do not need to explicitly compute  $\mathbf{w}$  and  $\phi(\mathbf{x})$  for kernel SVM.



- One common kernel: Radial Basis Function (RBF)

$$k(\mathbf{x}_i, \mathbf{x}) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}\|)$$

Gamma is for determining how the distance is considered influential.

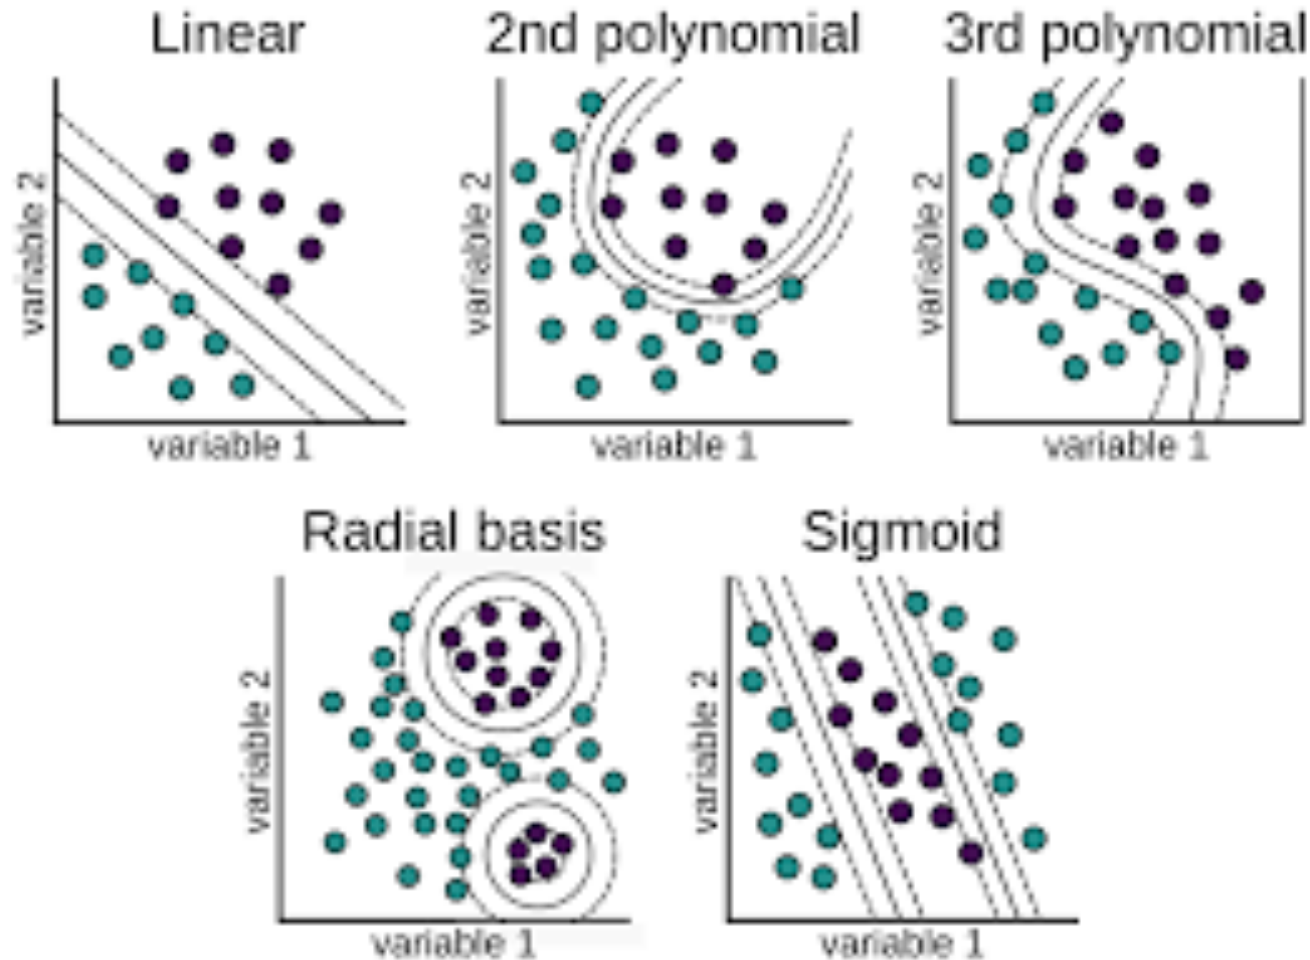


# SVM with RBF kernel



The learned decision boundary by SVM with RBF kernel is nonlinear in the original space

# SVM with Different Kernels and Decision Boundaries



# Outline for Data Preprocessing and Data Mining

- Data Preprocessing
- **Supervised learning**

- ❖ Regression

1. Linear regression with one variable
2. Linear Regression with multiple variables
3. The relationship between Correlation and Regression

- ❖ **Classification**


1. Perceptron
2. Artificial Neural Network
3. Support Vector Machine
4. **K Nearest Neighbor**

- **Unsupervised learning**

1. K-means Clustering
2. Hierarchical Clustering

# $k$ -NN Algorithm

- $k$  Nearest Neighbor Algorithm
  - $k$  is a user specified parameter, which means the number of nearest neighbors
- A Lazy Learning Algorithm
  - **Training:**
    - **No training process**, just store all training data in memory
  - **Prediction:**
    - Classify new samples based on most similar training samples via majority vote




- $x_q$  is the test sample.
- Assume  $k$  is equal to 5.
- Three out of its 5 nearest neighbors are from negative class.
- The predicted label for  $x_q$  is negative.

# Nearest Neighbors

- Training Data  $\{\mathbf{X}, \mathbf{y}\}$
- A test data point:  $\mathbf{x}_{\text{test}}$
- Idea: the label of a test data point is estimated from the known label(s) of the nearest neighbors of  $\mathbf{x}_{\text{test}}$  in the training data.
- Euclidean distance between feature vectors can be used to decide the nearest neighbors

$$\|\mathbf{x}_i - \mathbf{x}_j\|_2 = \sqrt{\sum_{k=1}^n (x_{i,k} - x_{j,k})^2}$$



# $k$ -NN algorithm

- Input: training data  $\{\mathbf{X}, y\}$ , a test data sample  $\mathbf{x}_{\text{test}}$ , parameter  $k$ 
  - Compute the distances between the test sample  $\mathbf{x}_{\text{test}}$  and each training data sample
  - Sort by distances and get the  $k$  nearest neighbors of  $\mathbf{x}_{\text{test}}$  ( $k$  is usually set to an odd number to prevent tie situations)
  - Use majority vote to predict the class label of  $\mathbf{x}_{\text{test}}$
- Output: predicted class label of  $\mathbf{x}_{\text{test}}$


# $k$ -Nearest Neighbors Example ( $k=3$ )

Training Data:

$x_1$	$x_2$	label
1	3	1
1	1	1
-1	-1	-1
-2	1	-1


Test Data:

$x_1$	$x_2$	label
-1	4	?



# $k$ -Nearest Neighbors Example ( $k=3$ )

$$\sqrt{(-1 - (-2))^2 + (4 - 1)^2} = \sqrt{10}$$



Training Data:

$x_1$	$x_2$	label
1	3	1
1	1	1
-1	-1	-1
-2	1	-1

Test Data:

$x_1$	$x_2$	label
-1	4	1

	$x_{\text{test}}$
$x_1$	$\sqrt{5}$
$x_2$	$\sqrt{13}$
$x_3$	$\sqrt{25}$
$x_4$	$\sqrt{10}$

Sort



	$x_{\text{test}}$
$x_1$	$\sqrt{5}$
$x_4$	$\sqrt{10}$
$x_2$	$\sqrt{13}$
$x_3$	$\sqrt{25}$




# Other $k$ -NN Distance metrics

## Minkowski Distance

$$d = \left( \sum_{i=1}^m |x_i - y_i|^p \right)^{1/p}$$

## Manhattan Distance

$$d = \sum_{i=1}^m |x_i - y_i|$$



Manhattan Distance

# Other $k$ -NN Distance metrics

Cosine Distance

$$\cos \theta = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \cdot \|\vec{b}\|}$$

Jaccard Distance

$$J(A,B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}$$

$A = \{X, Y, Z\}$  ,  $B = \{W, X\}$ ,  $C = \{X, Y\}$

$J(A,B) = 1 / 4$

$J(A,C) = 2 / 3$

# $k$ -NN - Generalize to multiple classes


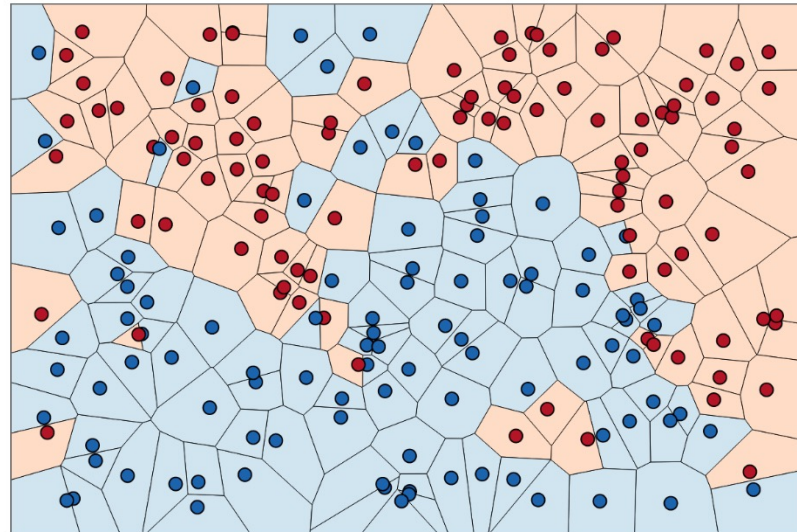



Image showing how similar data points typically exist close to each other

# $k$ -NN: Decision Boundaries

- $k$ -NN algorithm does not explicitly compute decision boundaries, but the decision boundaries can be inferred.
- Decision boundaries of 1-NN: Voronoi diagram
  - Show how input space divided into classes
  - Each line segment is equidistant between two neighboring data points.




# The Effect of $k$ in $k$ -NN



- If  $k$  is too small,  $k$ -NN will be very sensitive to “noisy samples” and lead to noisy decision boundaries.
- Large  $k$  will smooth the decision boundaries and may lead to better performance.

# The Effect of $k$ in $k$ -NN

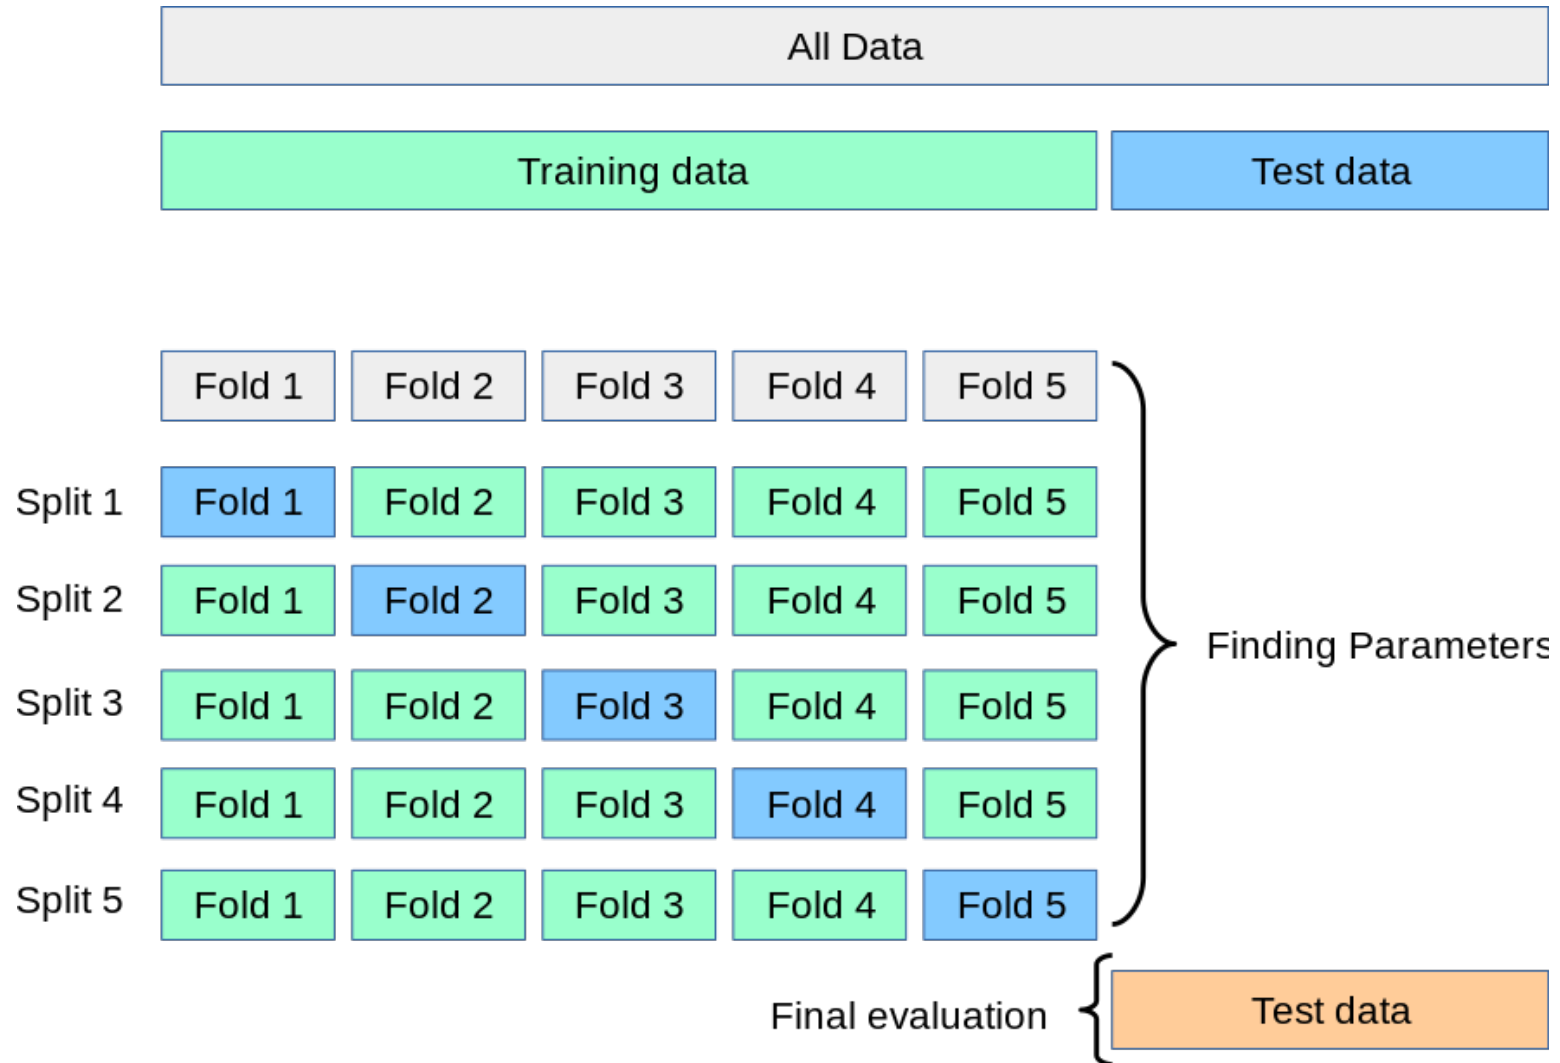


- If  $k$  is smaller than 5, data sample  $x$  is correctly predicted as red class.
- For larger  $k$ , data sample  $x$  is wrongly predicted as blue class.
- Therefore, if  $k$  is too large, we may end up with over-smoothed boundaries since it looks neighbors that are far away from the test data sample.

# How to choose $k$ ?

- $k$  being too small will be very sensitive to “noisy samples”, leading to noisy decision boundaries.
- $k$  being too large will lead to over-smoothed boundaries since it looks neighbors that are far away from the test data sample.
- We can use cross validation to find  $k$ .
  - Try several values of  $k$  : {5, 10, 20, 30}
  - Select the best  $k$  based on cross validation.

# Cross validation





# $k$ -NN: Some Issues and Remedies

- If some features (columns of data matrix) have large ranges, they will dominate the calculation of the distance.
  - Data Normalization
    - Min-max normalization: scale the range of each feature to be in range  $[0, 1]$
    - Decimal normalization: scale each feature to be in range  $(-1, 1)$
    - Z-score normalization: scale each feature to follow standard normal distribution
- Irrelevant, correlated features add noise to distance measures
  - Eliminate irrelevant features
  - Principal Component Analysis to reduce correlation among features
- Categorical Features, e.g., {'red', 'green', 'blue'}
  - One-hot encoding